
Logistic Regression with L_2 Regularization

Adrian Guthals, David Larson, Jason Yao
CSE 250B: Project #1
University of California, San Diego

1 INTRODUCTION

This paper will attempt to recreate the results by [1].

Highlight what Logistic Regression, why is it used and where is it used (real-world applications). Also mention potential pitfalls of the method (e.g. overfitting).

2 DESIGN AND ANALYSIS OF ALGORITHMS

Stochastic gradient descent (SGD) and Limited-memory Broyden-Fletcher-Goldfarb-Shannon (L-BFGS) were each implemented in Matlab to maximize the total conditional log likelihood of each training data set. In brief, SGD incorporated a fixed learning rate λ to control the change in log likelihood, which was averaged over random mini-batches of size κ . To control over-fitting, logistic regression was used and change in the objective function was evaluated on a separate validation dataset containing 30% of all training examples selected at random. Convergence was reached when change in the objective function was less than ω or the total number of epochs was greater than ε (which ever came first). (Add brief overview of L-BFGS and cite MinFunc's implementation)

minFunc[2]

For each of these algorithms, the input training data is formatted as a set of n examples $x_1 \dots x_n$ where each x_j is a real-valued vector of d features. Each x_i is correlated to a binary (Bernoulli) outcome y_i by a global parameter vector β of length $d + 1$. We assume this correlation follows the model below where $x_{i0} = 1$ for all i .

$$p_i = p(y_i|x_i; \beta) = \frac{1}{1 + \exp - (\sum_{j=0}^d \beta_j x_{ij})} \quad (1)$$

2.1 L2 REGULARIZATION

In order to prevent overfitting of machine learning, a penalty was imposed to regulate the values of the parameters. A regularization constant μ was introduced into the LCL objective function:

$$\hat{B} = \operatorname{argmax}_{\beta} LCL - \mu \|\beta\|_2^2 \quad (2)$$

where $\|\beta\|_2^2$ is the L_2 norm of the parameter vector. With this revision the derivative of the LCL becomes:

$$\frac{\partial}{\partial \beta_j} [\log p(y|x; \beta) - \mu \sum_{j=0}^d \beta_j^2] = (y - p)x_j - 2\mu\beta_j \quad (3)$$

2.2 STOCHASTIC GRADIENT DESCENT

Our SGD implementation first randomized the order of input examples to avoid repeated computation of random numbers and partitioned the input data into $x_1 \dots x_k$ *training* examples and $x_{k+1} \dots x_n$ *validation* examples. Then sequential mini-batches of size $\kappa < k$ taken from the training set were used to update the parameter vector β (initialized to all zero values) by the following equation. The constant μ quantifies the trade-off between maximizing likelihood and minimizing parameter values for L_2 Regularization.

$$\beta := \beta + \frac{\lambda}{\kappa} [-2\mu\beta + \sum_{i=1}^{\kappa} (y_i - p_i) x_i] \quad (4)$$

After each update of β , absolute change in the objective $\hat{\beta}$ was computed over all validation examples with the following function.

$$\hat{\beta} = \mu \|\beta\|_2 + \sum_{i=k+1}^n -\log(p_i^{y_i} (1 - p_i)^{1-y_i}) \quad (5)$$

Convergence was reached when change in the objective reached a value less than ω . Convergence was also reached if the total number of epochs was greater than ε . With this configuration, the time to run the SGD is $O(nd)$.

2.2.1 Cross-Validation

SGD with cross-validation is represented by the following pseudo-code:

1. Initialize all parameter values β_j
2. Initialize all conditional probabilities p_i
3. Initialize the objective function and its difference value
4. Randomly divide the example set into two groups: testing set of m rows and validation set of $(n-m)$ rows
5. while(objective function difference \geq threshold AND number of epochs $<$ maximum epochs allowed
 - (a) Randomly pick a sample of s rows out of the testing set
 - (b) Update β_j for all features
 - (c) Update the objective function

The algorithm was run ten times to generate 10 vectors of parameters. The ten vectors of parameters were averaged to calculate one vector with cross-validation. The result cross-validated vector was used to calculate the test error and its variance over all the instances. The test error was defined as the average value of the loss function:

$$\text{test error} = \sum_{i=1}^n -\log(y_i|x_i; \beta) \quad (6)$$

2.3 LIMITED-MEMORY BFGS

Limited-memory Broyden-Fletcher-Goldfarb-Shannon (L-BFGS) is a quasi-Newton optimization method used to find local extrema.

3 DESIGN OF EXPERIMENTS

3.1 PRE-PROCESSING TRAINING DATA

normalization, concatation?

3.1.1 USPS-N

Digit 9 vs other digits

3.1.2 Web

anything special?

Table 1: Hyperparameters.

	λ	μ
SGD	1.0	1.0
SGD	0.1	1.0
SGD	0.01	1.0

3.2 HYPERPARAMETERS

3.2.1 Learning Rate

How did we determine λ (the learning rate).

3.2.2 Regularization

How did we determine μ (the regularization constant).

4 RESULTS OF EXPERIMENTS

Results of experiments. Comparison of methods and data sets.

4.1 USPS-N

How did SGD and L-BFGS perform on the USPS-N data sets.

4.2 WEB

How did SGD and L-BFGS perform on the Web data sets.

4.2.1 Figures

Figure 1: Comparison of SGD and L-BFGS for USPS-N (left) and Web (right) data sets.

5 FINDINGS AND LESSONS LEARNED

Findings and lessons learned.

References

- [1] N. Ding and S. Vishwanathan, “t-logistic regression,” in *Advances in Neural Information Processing Systems 23*, J. Lafferty, C. K. I. Williams, J. Shawe-Taylor, R. Zemel, and A. Culotta, Eds., 2010, pp. 514–522.
- [2] M. Schmidt. [Online]. Available: <http://www.di.ens.fr/~mschmidt/>